Chapter 2 (Part 1): Bayesian Decision Theory (Sections 2.1-2.2)

An Example

- “Sorting incoming fish on a conveyor according to species using optical sensing”

Species

Sea bass

Salmon

Let us build a patter recognition system that classifies between Sea Bass and Salmon

Fish Classification: Salmon vs. Sea Bass

- Set up a camera and take some sample images
- Preprocessing involves image enhancement and segmentation;
  - separate touching or occluding fishes and
  - extract fish contour

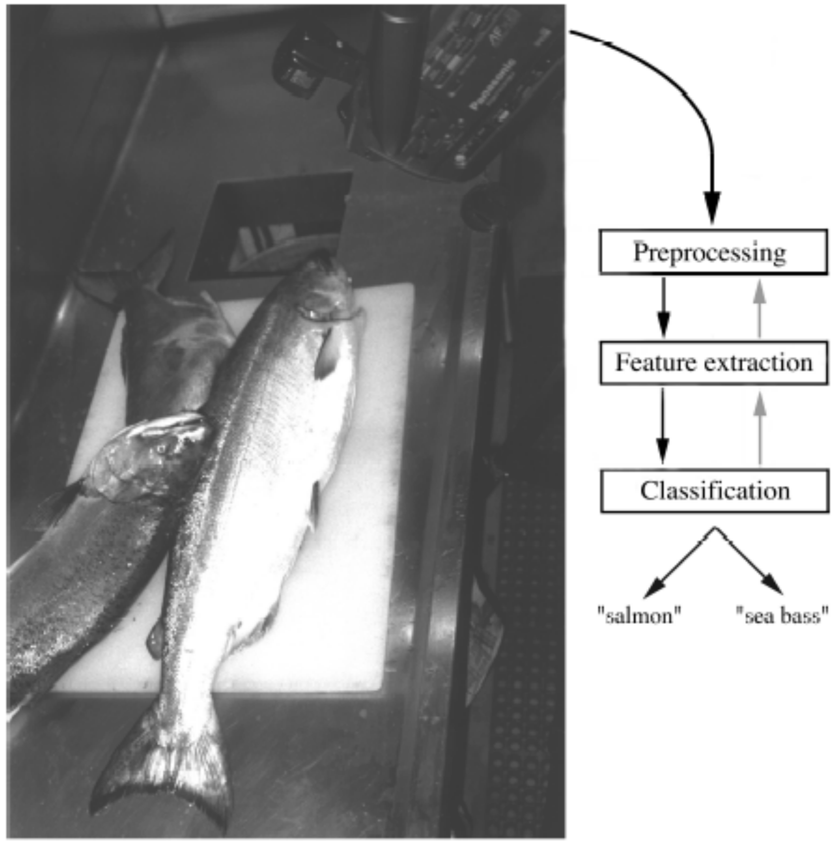

FIGURE 1.1. The objects to be classified are first sensed by a transducer (camera), whose signals are preprocessed. Next the features are extracted and finally the classification is emitted, here either "salmon" or "sea bass." Although the information flow is often chosen to be from the source to the classifier, some systems employ information flow in which earlier levels of processing can be altered based on the tentative or preliminary response in later levels (gray arrows). Yet others combine two or more stages into a unified step, such as simultaneous segmentation and feature extraction. From: Richard O. Duda, Peter E. Hart, and David G. Stork, *Pattern Classification*. Copyright © 2001 by John Wiley & Sons, Inc.

Problem Analysis

- Extract features from the images
  - Length
  - Lightness
  - Width
  - Number and shape of fins
  - Position of the mouth, etc…
- This is the set of all suggested features to explore for use in our classifier

Representation: Fish Length as Feature

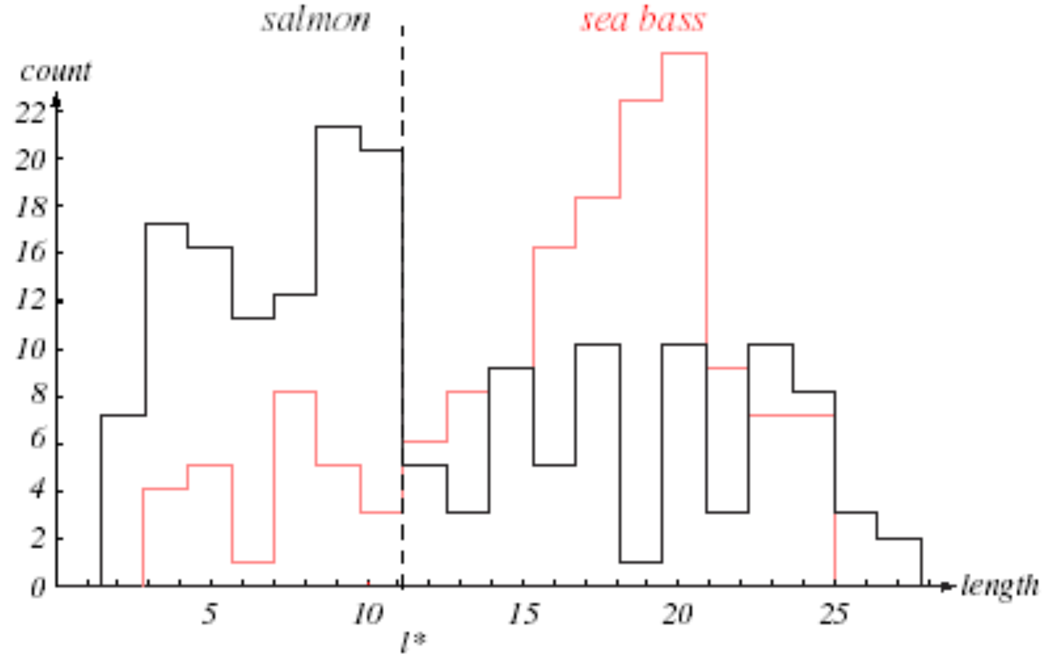

FIGURE 1.2. Histograms for the length feature for the two categories. No single threshold value of the length will serve to unambiguously discriminate between the two categories; using length alone, we will have some errors. The value marked l^* will lead to the smallest number of errors, on average. From: Richard O. Duda, Peter E. Hart, and David G. Stork, *Pattern Classification*. Copyright © 2001 by John Wiley & Sons, Inc.

Training (design or learning) Samples

State of Nature/Prior

- Prior probabilities reflect domain expert's knowledge of *how likely it is that each type of fish will appear,* before we actually see it.

  - State of nature is a random variable
  - $P(\omega_1) = P(\omega_2)$ (uniform priors): The catch of salmon and sea bass is equi-probable
  - $P(\omega_1) + P(\omega_2) = 1$ (exclusivity and exhaustivity)

State of Nature/Prior…

- Decision rule with only the prior information
  - Decide $\omega 1$ if $P(\omega 1) > P(\omega 2)$ otherwise decide $\omega 2$
- Probability of correct classification
  $= max\{P(w1), P(w2)\}$
- Probability of error for this decision
  - $P(error) = min\{P(w1), P(w2)\}$

Class-conditional Probabilities

- Use of the class–conditional information
- $P(x \mid \omega 1)$ and $P(x \mid \omega 2)$ describe the difference in lightness between the populations of sea-bass and salmon

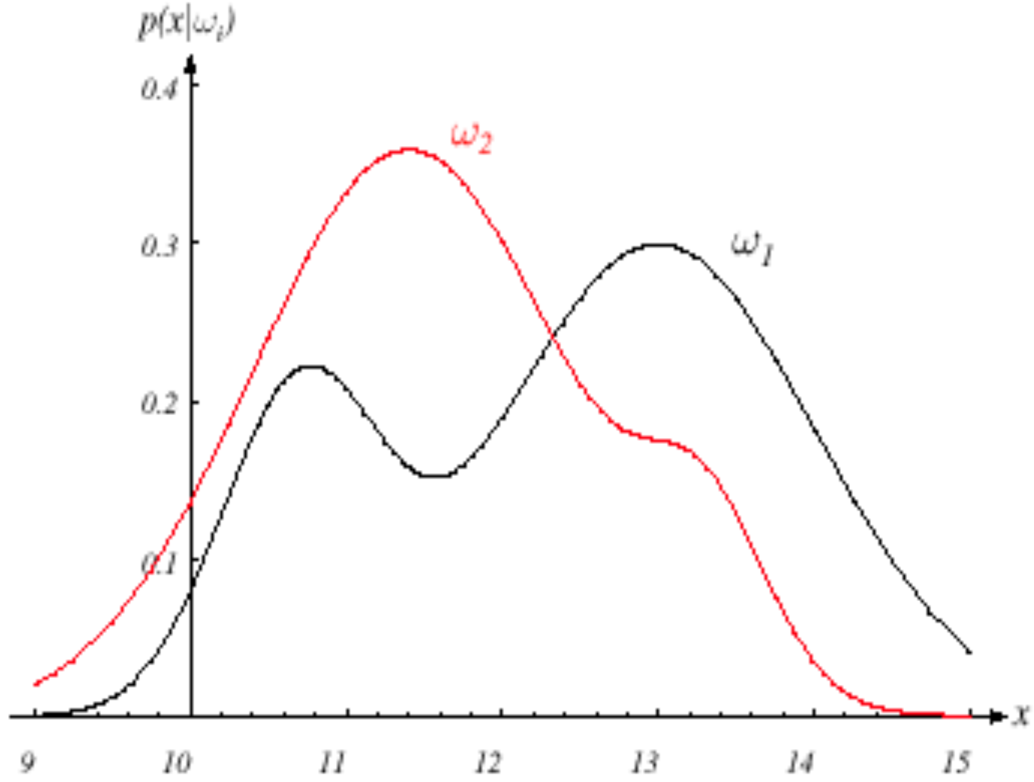

FIGURE 2.1. Hypothetical class-conditional probability density functions show the probability density of measuring a particular feature value x given the pattern is in category ω_i. If x represents the lightness of a fish, the two curves might describe the difference in lightness of populations of two types of fish. Density functions are normalized, and thus the area under each curve is 1.0. From: Richard O. Duda, Peter E. Hart, and David G. Stork, *Pattern Classification*. Copyright © 2001 by John Wiley & Sons, Inc.

Bayes' Classification

- Posterior, likelihood, evidence
  - P(ωj | x) = P(x | ωj) . P (ωj) / P(x)
  - In case of two categories

$$P(x) = \sum_{j=1}^{j=2} P(x \mid \omega_j) P(\omega_j)$$

  - Posterior = (Likelihood * Prior) / Evidence

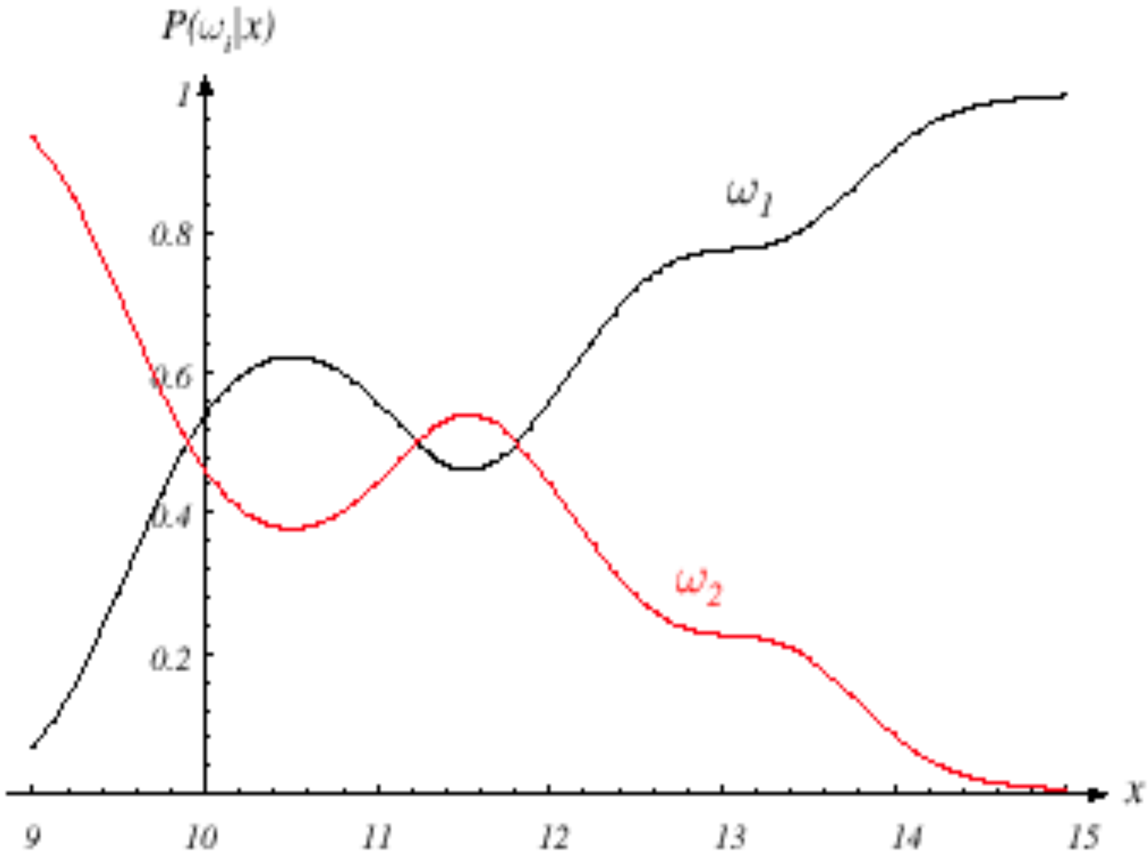

FIGURE 2.2. Posterior probabilities for the particular priors $P(\omega_1) = 2/3$ and $P(\omega_2) = 1/3$ for the class-conditional probability densities shown in Fig. 2.1. Thus in this case, given that a pattern is measured to have feature value $x = 14$, the probability it is in category ω_2 is roughly 0.08, and that it is in ω_1 is 0.92. At every x, the posteriors sum to 1.0. From: Richard O. Duda, Peter E. Hart, and David G. Stork, *Pattern Classification*.

Bayes' Decision

- Decision given the posterior probabilities
- X is an observation for which:
  - if $P(\omega 1 \mid x) > P(\omega 2 \mid x)$ True state of nature = $\omega 1$
  - if $P(\omega 1 \mid x) < P(\omega 2 \mid x)$ True state of nature = $\omega 2$
- Therefore, whenever we observe a particular x, the probability of error is :
  - $P(error \mid x) = P(\omega 1 \mid x)$ if we decide $\omega 2$
  - $P(error \mid x) = P(\omega 2 \mid x)$ if we decide $\omega 1$

Bayes' Decision

- Minimizing the probability of error

  Decide $\omega 1$ if $P(\omega 1 \mid x) > P(\omega 2 \mid x)$; otherwise decide $\omega 2$

- Therefore: $P(error \mid x) = \min [P(\omega 1 \mid x), P(\omega 2 \mid x)]$

Review

- Classification based on a single feature
- Two class classification
- Sample is assigned to one of the two classes
- The cost of making a false accept or a false reject is same

Bayesian Decision Theory

- Generalization of the preceding ideas
  - Use of more than one feature
  - Use more than two states of nature
  - Allowing actions other than decide on the state of nature
    - Allowing actions other than classification primarily allows the possibility of rejection
    - Refusing to make a decision in close or bad cases!
  - Introduce a loss function which is more general than the probability of error
    - The loss function states how costly each action taken is

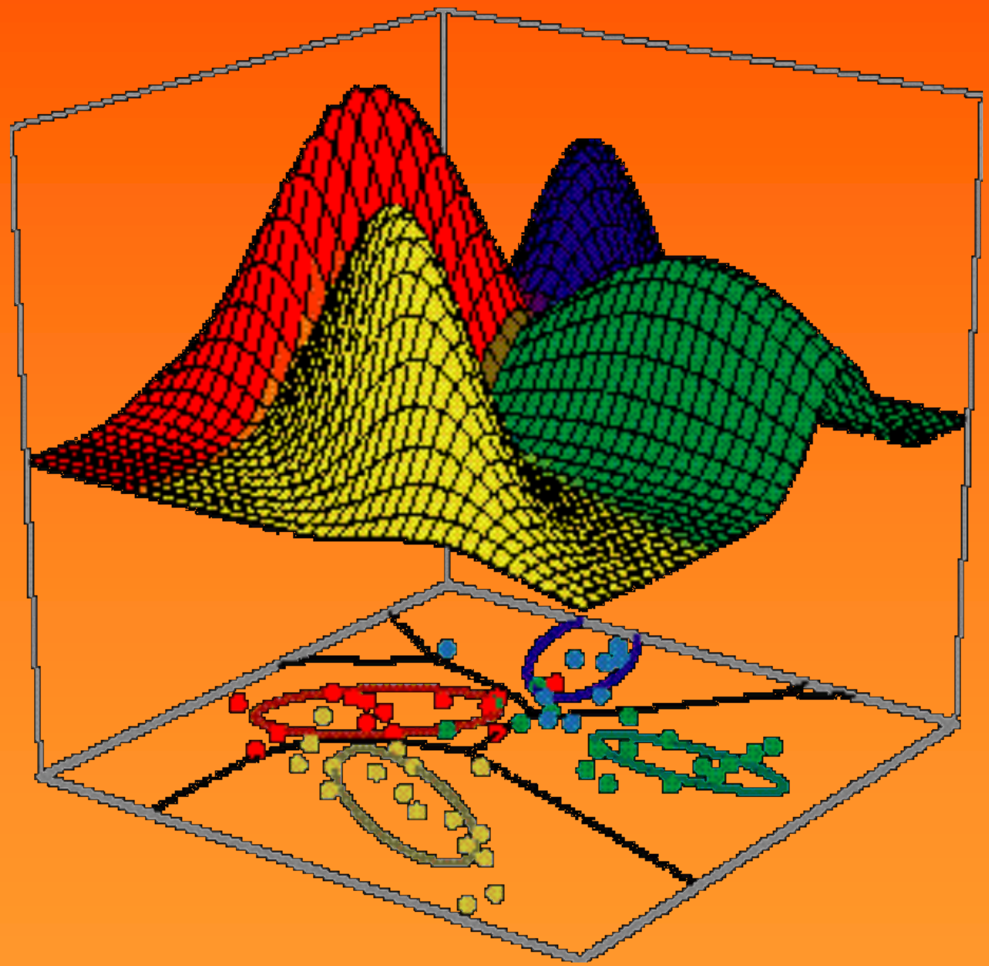

Pattern Classification

All materials in these slides were taken from Pattern Classification (2nd ed) by R. O. Duda, P. E. Hart and D. G. Stork, John Wiley & Sons, 2000 with the permission of the authors and the publisher